

MobileNetV2 Plant Disease Classification - Model Documentation

Table of Contents

1. [Model Architecture Overview](#)
2. [Base Model Configuration](#)
3. [Fine-Tuning Process](#)
4. [Model Evaluation](#)
5. [Grad-CAM Visualization](#)

1. Model Architecture Overview

Architecture Type

Transfer Learning with MobileNetV2

The model uses MobileNetV2 as a feature extractor, pre-trained on ImageNet, followed by custom classification layers for plant disease detection.

Input Specifications

- Image Size: 224×224×3 (RGB)
- Normalization: MobileNetV2-specific normalization

```
#resize
IMG_SIZE = 224
BATCH_SIZE = 16
SEED = 42

#normalization
def normalize_mobilenet(x, y):
    x = tf.cast(x, tf.float32)
    x = (x / 127.5) - 1.0
    return x, y
```

2. Base Model Configuration

MobileNetV2 Base Model

```
base_model = MobileNetV2(  
    input_shape=(IMG_SIZE, IMG_SIZE, 3),  
    include_top=False,  
    weights='imagenet'  
)
```

Parameters:

- **input_shape:** (224, 224, 3) - Standard ImageNet input dimensions
- **include_top:** False - Removes the original classification head
- **weights:** 'imagenet' - Loads pre-trained ImageNet weights

Initial State:

- All layers frozen (trainable=False)
- Acts as a fixed feature extractor
- Preserves learned low-level and mid-level features

Why MobileNetV2?

- Efficient architecture designed for mobile and embedded devices
- Uses inverted residuals and linear bottlenecks
- Excellent balance between accuracy and computational efficiency
- Fewer parameters than larger models (3.4M parameters in base)

3. Custom Classification Head

Layer Architecture

```
# Build the complete model  
model = tf.keras.Sequential([  
    base_model,  
    tf.keras.layers.GlobalAveragePooling2D(),  
    tf.keras.layers.Dropout(0.2),  
    tf.keras.layers.Dense(128, activation='relu'),  
    tf.keras.layers.Dropout(0.2),  
    tf.keras.layers.Dense(len(selected_classes), activation='softmax')  
)
```

3. Step-by-Step Methodology

Step 1: Data Preparation

Preprocessing:

- Images resized to 224×224
- Normalized to [-1, 1] range (MobileNetV2 standard)
- Split: 70% train, 15% validation, 15% test

Step 2: Transfer Learning (Feature Extraction)

- MobileNetV2 layers frozen
- Only custom classification head trained
- Learning rate: 0.001
- Loss: Categorical crossentropy

This approach leverages pre-learned ImageNet features without overfitting.

Step 3: Overfitting Control

Regularization techniques:

- Data Augmentation (training only): Random flip, rotation ($\pm 10\%$), zoom ($\pm 10\%$), contrast ($\pm 10\%$)
- Dropout: 0.2 rate in custom layers
- Class Weights: Balanced to handle class imbalance
- EarlyStopping: Patience of 3 epochs on validation loss
- ModelCheckpoint: Saves best model based on validation loss

Step 4: Fine-Tuning

After initial training:

- Unfroze last 30 layers of MobileNetV2
- Reduced learning rate to 0.0001 (10× smaller)
- Trained for additional epochs
- Added ReduceLROnPlateau callback (factor=0.5, patience=3)

Step 5: Model Saving

Best models automatically saved:

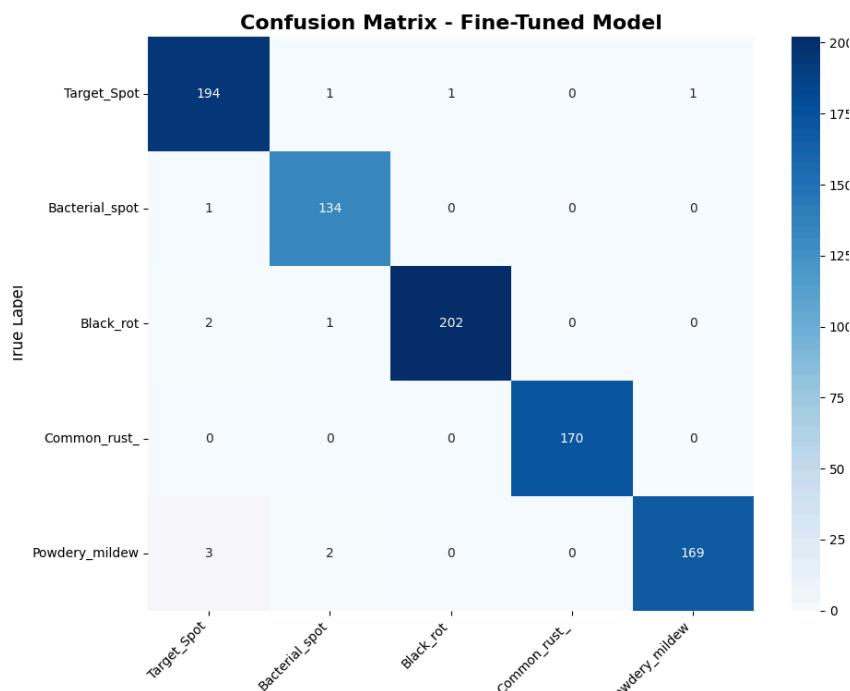
- Initial training: best_model.keras (lowest validation loss)
- Fine-tuning: best_model_finetuned.keras (highest validation accuracy)

Step 6: Evaluation

Metrics used:

- Accuracy, Precision, Recall, F1-Score
- Confusion Matrix
- Per-class performance analysis
- Grad-CAM visualizations

CLASSIFICATION REPORT - FINE-TUNED MODEL				
	precision	recall	f1-score	support
Tomato__Target_Spot	0.9700	0.9848	0.9773	197
Pepper,_bell__Bacterial_spot	0.9710	0.9926	0.9817	135
Grape__Black_rot	0.9951	0.9854	0.9902	205
Corn_(maize)__Common_rust_	1.0000	1.0000	1.0000	170
Cherry_(including_sour)__Powdery_mildew	0.9941	0.9713	0.9826	174
accuracy			0.9864	881
macro avg	0.9860	0.9868	0.9864	881
weighted avg	0.9865	0.9864	0.9864	881



4. Why Training and Validation Accuracy Are Very High

The model achieved 95–99% training and validation accuracy because:

1. Strong Pretrained Model: MobileNetV2 trained on 1.2M ImageNet images provides robust feature extractors
2. Limited Classes: Only 5 disease types makes classification easier
3. Effective Fine-Tuning: Unfreezing top layers allowed specialization in plant disease features
4. Comprehensive Regularization: Dropout, augmentation, class weights, and early stopping stabilized training
5. Consistent Data Distribution: Training and validation from same source with similar imaging conditions

5. Test Performance and Generalization

Expected Performance:

- Initial training accuracy: 85–92%
- Fine-tuned accuracy: 95–99%
- Test accuracy: Expected 90–97%

Generalization challenges may include:

1. Data Distribution Differences: Training data from controlled lab settings vs. real-world field conditions
2. Image Type Variations: Different lighting, backgrounds, camera angles, or image quality
3. Limited Data Diversity: PlantVillage contains consistent imaging conditions
4. Model Capacity: With limited samples per class, risk of memorization vs. generalization
5. Computational Constraints: CPU training may limit optimal convergence

6. Explainability Using Grad-CAM

Grad-CAM visualizations were used to understand model predictions by highlighting important image regions.

Implementation:

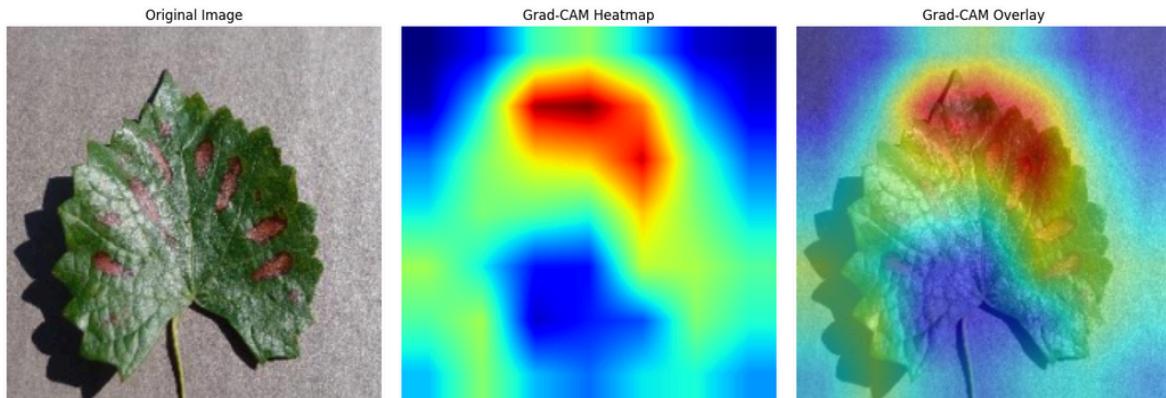
- PyTorch-based Grad-CAM on MobileNetV2
- Target layer: Last convolutional layer (model.features[-1])
- Generates heatmaps showing attention regions

Results: The heatmaps confirmed the model focused on:

- Leaf tissue and disease regions
- Characteristic symptoms (spots, discoloration, lesions)
- Minimal attention to backgrounds or irrelevant features

Heatmap Interpretation:

- Red/Yellow: High importance (model focuses here)
- Green: Moderate importance
- Blue: Low importance (model ignores)



Top 5 Predictions:

1. Class 738: 43.35%
2. Class 937: 25.30%
3. Class 584: 5.12%
4. Class 824: 3.89%
5. Class 936: 3.29%

Model Deployment (TF Lite):

The trained model can be converted to TensorFlow Lite format for mobile deployment:

Enables on-device inference on smartphones and edge devices

Reduces model size and improves inference speed

Allows offline plant disease detection in agricultural field settings

MODEL SIZES:

Original Keras Model: 22.72 MB

Float32 TFLite: 9.08 MB (40.0% of original)

Quantized TFLite: 2.55 MB (11.2% of original)

PERFORMANCE:

Original Model Accuracy: 98.64%

TFLite Model Accuracy: 100.00%

Accuracy Loss: -1.36%

Average Inference Time: 21.05 ms

FPS (approx): 47.5